

Lecture 4: Further solution concepts

Motivations

- We reason about multiplayer games using solution concepts, principles according to which we identify interesting subsets of the outcomes of a game
- Nash equilibrium is the most important solution concept
- There are also a large number of others:
 - Maximin and minmax strategies
 - Minimax regret
 - Removal of dominated strategies
 - Rationalizability
 - Correlated equilibrium
 - Trembling-hand perfect equilibrium

Maxmin and minmax strategies

Definition (Maxmin)

The maxmin strategy for player i is $s_i^* = \arg \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$ and the maxmin value for player i is $\max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$

- The **maxmin strategy** of player i in an n -players game is a strategy that maximizes i 's **worst – case payoff**, in the situation where all the others players happen to play the strategies which cause the greatest harm to i
- The **maxmin strategy** is a sensible choice for a **conservative agent** who wants to maximize his expected utility **without having to make any assumptions about the other agents**
- The **maxmin value** (or security level) of the game for player i is that minimum amount of payoff guaranteed by a **maxmin strategy**
- It is strategy that **defends against** other agents (defensive strategy)
- Player i set the mixed strategy \Rightarrow player $-i$ observe this strategy (not an action) and choose their own strategies to minimize i 's expected payoff

(temporal interpretation)

Maxmin and minmax strategies

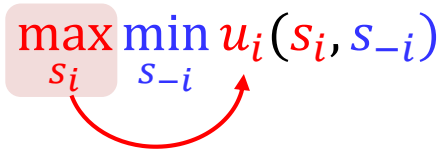
Definition (Minmax, two-player)

In an two-player game, the *minmax strategy* for player i against player $-i$ is $s_i^* = \arg \min_{s_i} \max_{s_{-i}} u_{-i}(s_i, s_{-i})$ and the minmax value is $\min_{s_i} \max_{s_{-i}} u_{-i}(s_i, s_{-i})$

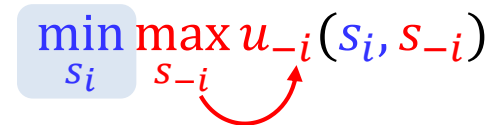
- The *minmax strategy* of player i in an two-players game is a strategy that keeps the maximum payoff of $-i$ at a minimum
- The *minmax value* of player $-i$ is that minimum
- It is strategy that *attack* against other agents (offensive strategy)

Maxmin and minmax strategies

In agent i 's perspective

$$\max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$$


- Agent always maximizes its payoff
- **Defensive strategy** (if max is first)

$$\min_{s_i} \max_{s_{-i}} u_{-i}(s_i, s_{-i})$$


- Agent always maximizes its payoff
- **offensive strategy** (if min is first)

Maxmin and minmax strategies

Definition (Minmax, n -player)

In an n -player game, the minmax strategy for **player i** against player $j \neq i$ is **i -th component** of the mixed-strategy profile s_{-j} in the expression $\arg \min_{s_{-j}} \max_{s_j} u_j(s_j, s_{-j})$. As before, the minmax value for player j is $\min_{s_{-j}} \max_{s_j} u_j(s_j, s_{-j})$

- Here, we assume that all the players other than j choose to “gang up” on j
 - They are able to coordinate appropriately when there is more than one strategy profile that would yield the same minimal payoff for j



Maxmin and minmax strategies

- An agent's maximum strategy nor his minmax strategy depends on the strategies that the other agents actually choose
- the maximin and minmax strategies give rise to solution concept
- Call $s = (s_1, \dots, s_n)$ a **maxmin strategy profile** of a given game if s_1 is a maxmin strategy for player 1, s_2 is a maxmin strategy for player 2 and so on.
 - Similar to **minmax strategy profile**
- In two-player games, a player's minmax value is always equal to his maxmin value

$$\min_{s_{-i}} \max_{s_i} u_i(s_i, s_{-i}) = \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$$

- For games with more than two players, a weaker condition holds:

$$\min_{s_{-i}} \max_{s_i} u_i(s_i, s_{-i}) \geq \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$$

- See that player $-i$ chooses first, allowing player i to best respond to it.

Minimax theorem (von Neumann, 1928)

Theorem (Minmax theorem by von Neumann, 1928)

In any finite, two-player, zero-sum game, in any **Nash equilibrium** each player receives a payoff that is equal to both his **maxmin** value and his **minmax** value.

- Minmax theorem states that in a two-player zero-sum game:

$$\text{maximin value} = \text{minmax value} = \text{Nash equilibrium value}$$

- Any *maximin* strategy profile is a Nash equilibrium. Furthermore, these are all the Nash equilibria
 - Consequently, all Nash equilibria have the same payoff vector

Minimax theorem (von Neumann, 1928)

Proof:

- Let's assume (s'_i, s'_{-i}) be an arbitrary Nash equilibrium and denote v_i to be the i 's equilibrium payoff
- Denote i 's maxmin value as $\bar{v}_i = \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i})$
- Denote i 's minmax value as $\underline{v} = \min_{s_{-i}} \max_{s_i} u_i(s_i, s_{-i})$
- First, we show that $\bar{v}_i = v_i$

$$\checkmark \quad \bar{v}_i \leq v_i$$

$$\bar{v}_i = \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i}) \leq \max_{s_i} u_i(s_i, s'_{-i}) < v_i$$

$$\checkmark \quad \bar{v}_i \geq v_i$$

$$v_{-i} = \max_{s_{-i}} u_{-i}(s'_i, s_{-i})$$
$$-v_{-i} = \min_{s_{-i}} -u_{-i}(s'_i, s_{-i})$$

since the game is zero sum, $-v_{-i} = v_i$ and $u_i = -u_{-i}$, thus

$$v_i = \min_{s_{-i}} u_i(s'_i, s_{-i})$$
$$\bar{v}_i = \max_{s_i} \min_{s_{-i}} u_i(s_i, s_{-i}) \geq \min_{s_{-i}} u_i(s'_i, s_{-i}) = v_i$$

➤ As a result, $v_i = \bar{v}_i$


Minimax theorem example

		Player 2	
		q Heads	$1 - q$ Tails
Player 1	p Heads	1, -1	-1, 1
	$1 - p$ Tails	-1, 1	1, -1

- Player 1's maxmin value :
$$\bar{u}_1 = \max_{s_1} \min_{s_2} u_1(s_1, s_2)$$
$$= \max_p \min_q \{pq - p(1 - p) - (1 - p)q + (1 - p)(1 - q)\}$$
- Player 1's minmax value :
$$\underline{u}_1 = \min_{s_2} \max_{s_1} u_1(s_1, s_2)$$
$$= \min_q \max_p \{pq - p(1 - p) - (1 - p)q + (1 - p)(1 - q)\}$$

Minimax theorem example

- Player 1's maxmin value :
$$\bar{u}_1 = \max_{s_1} \min_{s_2} u_1(s_1, s_2)$$
$$= \max_p \min_q \{pq - p(1-p) - (1-p)q + (1-p)(1-q)\}$$
- For any p set by player 1, player 2 tries to choose q **deterministically** to minimize u_1
 - When player 2 plays Heads ($q = 1$): $u_1(p, H) = 2p - 1$
 - When player 2 plays Tails ($q = 0$): $u_1(p, T) = 1 - 2p$
- $\min_q \{pq - p(1-p) - (1-p)q + (1-p)(1-q)\} = \min\{2p - 1, 1 - 2p\}$
Selecting among the two possible alternatives
- Thus, $\bar{u}_1 = \max_p \min\{2p - 1, 1 - 2p\} = 0$



- Player 1's maxmin strategy:


$$\bar{s}_1 = \operatorname{argmax}_{s_1} \min_{s_2} u_1(s_1, s_2) = \left(\frac{1}{2}, \frac{1}{2}\right)$$

- Player 1's maxmin value:

$$\bar{u}_1 = \max_{s_1} \min_{s_2} u_1(s_1, s_2) = 0$$

Minimax theorem example

- Player 1's minmax value :
$$\underline{u}_1 = \min_{s_2} \max_{s_1} u_1(s_1, s_2)$$
$$= \min_q \max_p \{pq - p(1-p) - (1-p)q + (1-p)(1-q)\}$$
- For any q set by player 2, player 1 tries to choose p deterministically to maximize u_1
 - When player 1 plays Heads ($p = 1$): $u_1(H, q) = 2q - 1$
 - When player 1 plays Tails ($p = 0$): $u_1(T, q) = 1 - 2q$
- $\max_p \{pq - p(1-p) - (1-p)q + (1-p)(1-q)\} = \max\{2q - 1, 1 - 2q\}$
- Thus, $\underline{u}_1 = \min_q \max\{2q - 1, 1 - 2q\} = 0$



- Player 2's minmax strategy:

$$\underline{s}_2 = \operatorname{argmin}_{s_2} \max_{s_1} u_1(s_1, s_2) = \left(\frac{1}{2}, \frac{1}{2}\right)$$


- Player 1's minmax value:

$$\underline{u}_1 = \min_{s_2} \max_{s_1} u_1(s_1, s_2) = 0$$

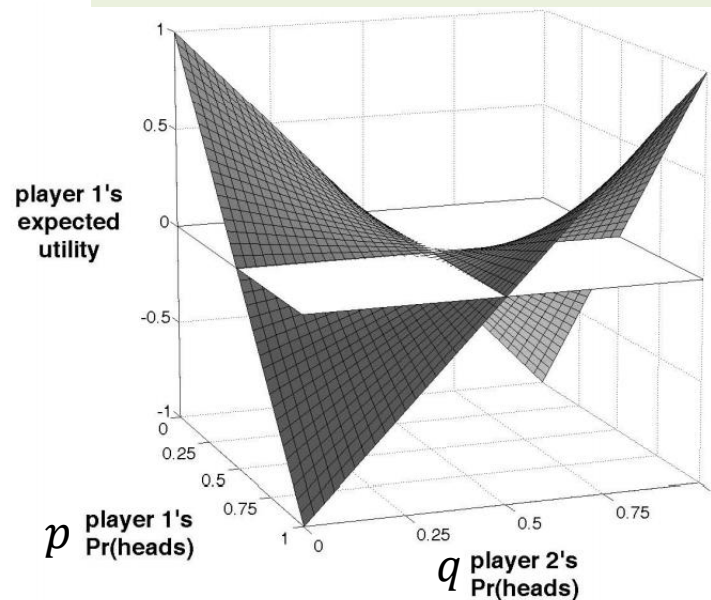
Minimax theorem graphical representation

		Player 2	
		q Heads	$1 - q$ Tails
Player 1	p Heads	1, -1	-1, 1
	$1 - p$ Tails	-1, 1	1, -1

$$u_1(p, q) = pq - p(1 - p) - (1 - p)q + (1 - p)(1 - q)$$





$$u_1(p, q) = pq - p(1 - p) - (1 - p)q + (1 - p)(1 - q)$$



- Nash equilibria in zero-sum games can be viewed graphically as a “saddle” in a high-dimensional space
- At a saddle point, any deviation of the agent lowers his utility and increases the utility of the other agent.

Minimax theorem graphical representation

Player 2's
new strategy



- Nash equilibria in zero-sum games can be viewed graphically as a “saddle” in a high-dimensional space
- At a saddle point, any deviation of the agent lowers his utility and increases the utility of the other agent.

Minimax regret

		Player 2	
		L	R
Player 1	T	100, a	$1 - \epsilon$, b
	B	2, c	1, d

- We argued agents might play maxmin strategies to achieve good payoffs in the worst case
- **Player 1's maximin strategy is to play B :**
 - If player 1 play T , then player 2 will chose R to minimize player 1's payoff: $u_1 = 1 - \epsilon$
 - If player 1 play B , then player 2 will chose R to minimize player 1's payoff: $u_1 = 1$
 - Thus, maximin strategy for player 1 is to play B , giving him a payoff of 1

Minimax regret

		Player 2	
		<i>L</i>	<i>R</i>
Player 1	<i>T</i>	100, <i>a</i>	$1 - \epsilon$, <i>b</i>
	<i>B</i>	2, <i>c</i>	1, <i>d</i>

- However, the other agent is **not believed to be malicious**, but is instead **unpredictable**
- In this case, agents might care about minimizing their worst-case losses, rather than maximizing their worst case payoffs
- **Player 1's Minmax regret strategy is to play *T*:**
 - If player 2 were to play *R*, then it would not matter very much how player 1 plays
 - ✓ The most he could lose by playing the wrong way would be ϵ
 - If player 2 were to play *L*, then player 1's action would be very significant
 - ✓ If player makes wrong choice, his utility would be decreased by 98
 - Thus, given that player can maximize your regret, player 1 might choose to play *T* in order to minimize his worst-case loss

Minimax regret

Definition (Regret)

An agent i 's regret for playing an action a_i if the other agents adopt action profile a_{-i} is defined as

$$\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i})$$

- In words, this is the amount that i loses by playing a_i , rather than playing his best response to a_{-i} . Of course, i does not know that actions the other players will take.
- But, we can consider those actions that would give him the highest regret for playing a_i

Minimax regret

Definition (Regret)

An agent i 's regret for playing an action a_i if the other agents adopt action profile a_{-i} is defined as

$$\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i})$$

- In words, this is the amount that i loses by playing a_i , rather than playing his best response to a_{-i} . Of course, i does not know that actions the other players will take.
- But, we can consider those actions that would give him the highest regret for playing a_i

Definition (Max Regret)

An agent i 's regret for playing an action a_i if the other agents adopt action profile a_{-i} is defined as

$$\max_{a_{-i} \in A_{-i}} \left(\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i}) \right)$$

- This is the amount that i loses by playing a_i rather than playing his best response to a_{-i} , if the other agents chose the a_{-i} that makes this loss as large as possible

Minimax regret

Definition (Regret)

An agent i 's regret for playing an action a_i if the other agents adopt action profile a_{-i} is defined as

$$\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i})$$

- In words, this is the amount that i loses by playing a_i , rather than playing his best response to a_{-i} . Of course, i does not know that actions the other players will take.
- But, we can consider those actions that would give him the highest regret for playing a_i

Definition (Max Regret)

An agent i 's regret for playing an action a_i if the other agents adopt action profile a_{-i} is defined as

$$\max_{a_{-i} \in A_{-i}} \left(\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i}) \right)$$

- This is the amount that i loses by playing a_i rather than playing his best response to a_{-i} , if the other agents chose the a_{-i} that makes this loss as large as possible

Definition (Minmax Regret)

Minmax regret actions for agent i are defined as

$$\operatorname{argmin}_{a_i \in A_i} \left[\max_{a_{-i} \in A_{-i}} \left(\left[\max_{a'_i \in A_i} u_i(a'_i, a_{-i}) \right] - u_i(a_i, a_{-i}) \right) \right]$$

- **Minmax regret actions are one** that yields the smallest maximum regret

Removal of dominated strategies

Definition (Domination)

Let s_i and s'_i be two strategies of player i , and S_{-i} the set of all strategy profiles of the remaining players. Then,

1. s_i **strictly dominates** s'_i if for all $s_{-i} \in S_{-i}$, it is the case that $u_i(s_i, s_{-i}) > u_i(s'_i, s_{-i})$
2. s_i **weekly dominates** s'_i if for all $s_{-i} \in S_{-i}$, it is the case that $u_i(s_i, s_{-i}) \geq u_i(s'_i, s_{-i})$, and for at least one $s_{-i} \in S_{-i}$, it is the case that $u_i(s_i, s_{-i}) > u_i(s'_i, s_{-i})$
3. s_i **very weekly dominates** s'_i if for all $s_{-i} \in S_{-i}$, it is the case that $u_i(s_i, s_{-i}) \geq u_i(s'_i, s_{-i})$

- Domination is comparison between two strategies s_i and s'_i given others $s_{-i} \in S_{-i}$

Definition (Pareto domination)

Strategy profile s Pareto dominates **strategy profile** s' if for all $i \in N$, $u_i(s) \geq u_i(s')$, and there exists some $j \in N$ for which $u_j(s) > u_j(s')$.

Removal of dominated strategies

Definition (Dominant strategy)

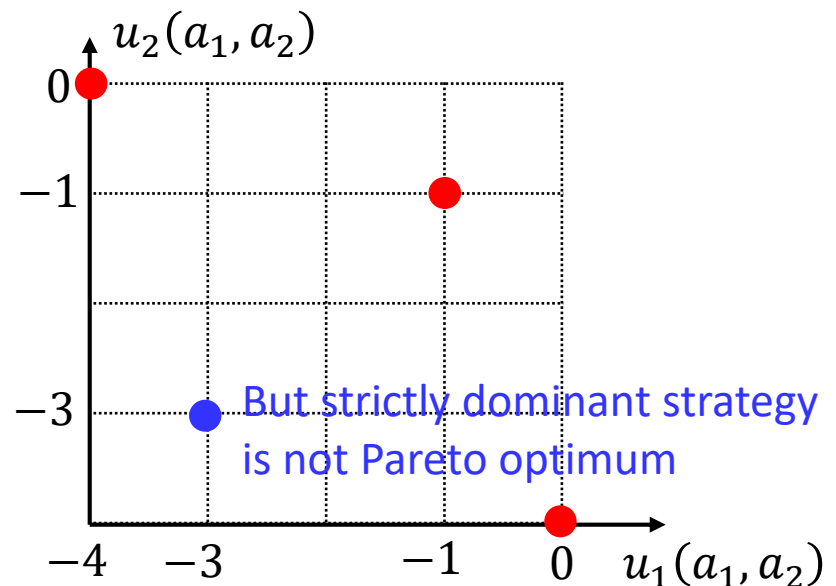
A strategy is strictly (resp., weakly; very weakly) dominant **for an agent** if it strictly (weakly; very weakly) dominates any other strategy for that agent.

- A strategy profile (s_1, \dots, s_n) in which every s_i is dominant for player i (whether strictly, weakly, or very weakly) is a Nash equilibrium.
- A strategy profile consisting of dominant strategies for every player must be a Nash equilibrium
 - An equilibrium in strictly dominant strategies must be unique.

	C	D
C	-1, -1	-4, 0
D	0, -4	-3, -3

Prisoner's Dilemma game

Strictly
dominant
strategy profile



Removal of dominated strategies

Definition (Dominated strategy)

A strategy s_i is strictly (weakly; very weakly) dominated for an agent i if some other strategy s'_i strictly (weakly; very weakly) dominates s_i

- Note that it is easy to check if a strategy is dominated because we need to find any strategy that dominate it

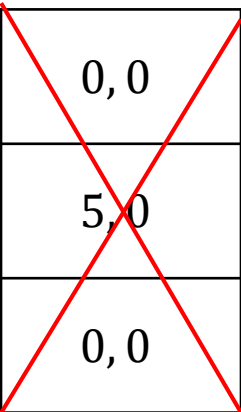
Removal of dominated strategies

	L	C	R
U	3, 1	0, 1	0, 0
M	1, 1	1, 1	5, 0
D	0, 1	4, 1	0, 0

- R is dominated by L

Removal of dominated strategies

	L	C	R
U	3, 1	0, 1	0, 0
M	1, 1	1, 1	5, 0
D	0, 1	4, 1	0, 0



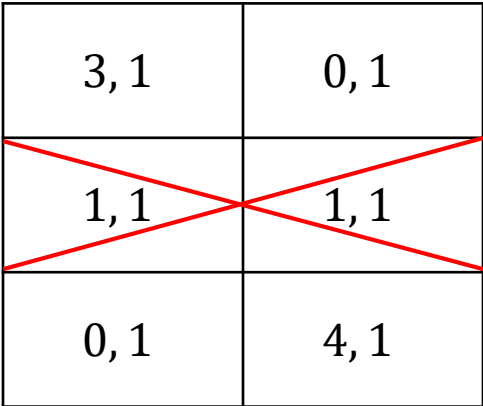
- R is dominated by L

Removal of dominated strategies

	L	C
U	3, 1	0, 1
M	1, 1	1, 1
D	0, 1	4, 1

Removal of dominated strategies

	L	C
U	3, 1	0, 1
M	1, 1	1, 1
D	0, 1	4, 1



- M is dominated by the mixed strategy that selects U and D with equal probability

Removal of dominated strategies

	L	C
U	3, 1	0, 1
D	0, 1	4, 1

- No other strategies are dominated.

Removal of dominated strategies

	L	C	R
U	4, 3	5, 1	6, 2
M	2, 1	8, 4	3, 6
D	3, 0	9, 6	2, 8

- Find an equilibrium by yourself

Removal of dominated strategies

- This process **preserves Nash equilibria**.
 - strict dominance: all equilibria preserved.
 - weak or very weak dominance: at least one equilibrium preserved.
- Thus, it can be used as a **preprocessing step** before computing an equilibrium
 - Some games are solvable using this technique.
 - Example: Prisoner's Dilemma!
- What about the **order of removal** when there are multiple dominated strategies?
 - strict dominance: doesn't matter.
 - weak or very weak dominance: can affect which equilibria are preserved.

	L	C
U	1, 1	2, 1
D	1, 2	3, 1

- Remove the action of the column player first
- Remove the action of the row player first

What is the result?

Removal of dominated strategies

Cournot duopoly

- Two identical firms, players 1 and 2, produce some good
- Firm i produce quantity q_i
- Cost for production is $c_i(q_i) = 10q_i$
- Price is given by $d = 100 - (q_1 + q_2)$
- The profit of company 1 is $u_1(q_1, q_2) = (100 - q_1 - q_2)q_1 - 10q_1 = 90q_1 - q_1^2 - q_1q_2$

What should firm 1 do in order to maximize their profit?

Removal of dominated strategies

Cournot duopoly

- Two identical firms, players 1 and 2, produce some good
- Firm i produce quantity q_i
- Cost for production is $c_i(q_i) = 10q_i$
- Price is given by $d = 100 - (q_1 + q_2)$
- The profit of company 1 is $u_1(q_1, q_2) = (100 - q_1 - q_2)q_1 - 10q_1 = 90q_1 - q_1^2 - q_1q_2$

What should firm 1 do in order to maximize their profit?

- As the payoff is concave in q_1 , the maximum is obtained by imposing the derivative of the payoff with respect q_1 for any given value of q_2

$$q_1 = \frac{90 - q_2}{2}$$

➤ That is, for any given q_2 chosen by company 2, company maximize its payoff

- The same applied to company 2

$$q_2 = \frac{90 - q_1}{2}$$

Removal of dominated strategies

Cournot duopoly

- The profit of company 1 is $u_1(q_1, q_2) = (100 - q_1 - q_2)q_1 - 10q_1 = 90q_1 - q_1^2 - q_1q_2$
- As the payoff is concave in q_1 , the maximum is obtained by imposing the derivative of the payoff with respect q_1 for any given value of q_2


$$q_1 = \frac{90 - q_2}{2}$$

- Company 1 will never choose to produce more than $q_1 > 45$ because any quantity $q_1 > 45$ is strictly dominated by $q_1 = 45$ as follows:
 - $u_1(q_1 = 45, q_2) = (100 - 45 - q_2)45 - 450 = 2025 - 45q_2$
 - $u_1(q_1, q_2) = (100 - q_1 - q_2)q_1 - 10q_1 = 90q_1 - q_1^2 - q_1q_2$
 - $u_1(45, q_2) - u_1(q_1, q_2) = 2025 - q_1(90 - q_1) - q_2(45 - q_1) > 0$
for any $q_1 > 45$ regardless of q_2
- Due to symmetry, any $q_2 > 45$ is strictly dominated by $q_2 = 45$
- **The first round of iterated elimination:**
 - A rational produces no more than 45 units, implying that the effective strategy space that survives one round of elimination is $q_i \in [0, 45]$ for $i \in \{1, 2\}$

Removal of dominated strategies

Cournot duopoly


- **The first round of iterated elimination:**
 - $q_2 > 45$ is strictly dominated by $q_2 \leq 45$



Removal of dominated strategies

Cournot duopoly


- **The second round of iterated elimination:**
 - Because $q_2 \leq 45$, the equation $q_1 = \frac{90 - q_2}{2}$ implies that company 1 will choose $q_1 \geq 22.5$
 - Symmetric argument applies to $q_2 \geq 22.5$
 - Therefore the second round of elimination implies that the surviving strategy sets are $q_i \in [22.5, 45]$ for $i \in \{1, 2\}$



Removal of dominated strategies

Cournot duopoly

- **The third round of iterated elimination:**
 - Because $q_2 \geq 22.5$, the equation $q_1 = \frac{90 - q_2}{2}$ implies that company 1 will choose $q_1 \leq 33.75$
 - Symmetric argument applies to $q_2 \leq 33.75$
 - Therefore the second round of elimination implies that the surviving strategy sets are $q_i \in [22.5, 33.75]$ for $i \in \{1, 2\}$



Rationalizability

- Rather than ask what is irrational, ask what is a best response to some beliefs about the opponent
 - assumes opponent is rational
 - assumes opponent knows that you and the others are rational
 - ...
- Examples
 - is heads rational in matching pennies?
 - is cooperate rational in prisoner's dilemma?
- Will there always exist a rationalizable strategy?
 - Yes, equilibrium strategies are always rationalizable.
- Furthermore, in two-player games, rationalizable \Leftrightarrow survives iterated removal of strictly dominated strategies.

If there is intelligent life on other planets, in a majority of them, they would have discovered correlated equilibrium before Nash equilibrium.

– Roger Myerson

Correlated equilibrium

- Consider again Battle of the Sexes.
 - The values of each player under mixed Nash equilibrium is $\frac{2}{3}$

		Player 2	
		p	$1 - p$
		TF	LA
Player 1	q	2, 1	0, 0
	$1 - q$	0, 0	1, 2
		TF	LA

$$\begin{aligned}
 u_1(\text{TF}) &= u_1(\text{LA}) \\
 2 \times p + 0 \times (1 - p) &= 0 \times p + 1 \times (1 - p) \\
 p &= \frac{1}{3}
 \end{aligned}$$

$$\begin{aligned}
 u_2(\text{TF}) &= u_2(\text{LA}) \\
 1 \times q + 0 \times (1 - q) &= 0 \times q + 2 \times (1 - q) \\
 q &= \frac{2}{3}
 \end{aligned}$$

- The mixed Nash equilibrium is $s^* = (s_1^*, s_2^*) = \left\{ \left(\frac{2}{3}, \frac{1}{3} \right), \left(\frac{1}{3}, \frac{2}{3} \right) \right\}$
- The expected payoff under s^* are $u_1^* = \frac{2}{3} = u_2^*$

Can we do better?

Correlated equilibrium

- Consider again Battle of the Sexes.
 - The values of each player under mixed Nash equilibrium is $2/3$

		Player 2	
		p	$1 - p$
		TF	LA
Player 1	q	TF	2, 1
	$1 - q$	LA	0, 0
			1, 2

- We could use the same idea to achieve the fair outcome in battle of the sexes.
 - Intuitively, the best outcome seems a 50-50 split between (TF, TF) and (LA, LA).

$$u_1^{CE} = \frac{1}{2}(2 + 1) = \frac{3}{2} > u_1^{NE} = \frac{2}{3}$$

$$u_2^{CE} = \frac{1}{2}(2 + 1) = \frac{3}{2} > u_2^{NE} = \frac{2}{3}$$

We show that **no player has an incentive to deviate** from the “recommendation” of the coin.

Correlated equilibrium

- Another classic example: traffic game

	Go	Wait
Go	-100, -100	10, 0
Wait	0, 10	-10, -10

Traffic game



- What is the natural solution here?
 - A traffic light: a fair randomizing device that tells one of the agents to go and the other to wait.
- Benefits:
 - the negative payoff outcomes are completely avoided
 - fairness is achieved
 - the sum of social welfare exceeds that of mixed Nash equilibrium

Correlated equilibrium

- More complex example

	L	R
U	5, 1	0, 0
D	4, 4	1, 5

- There are two pure strategy Nash equilibria: (U, L) and (D, R).
- This implies that there is a unique mixed strategy equilibrium with expected payoff (5/2, 5/2).
- Suppose the players find a mediator who chooses $x \in \{1, 2, 3\}$ with equal probability 1/3. She then sends the following messages:
 - If $x = 1$, player 1 plays U, player 2 plays L.
 - If $x = 2$, player 1 plays D, player 2 plays L.
 - If $x = 3$, player 1 plays D, player 2 plays R.

Actions are correlated
- Our first example presumed that everyone perfectly observes the random event; not required.
- More generally, some random variable with a commonly known distribution, and a private signal to each player about the outcome.
 - signal doesn't determine the outcome or others' signals; however, correlated:
✓ Actions for agents are jointly determined by a drawn random variable

Correlated equilibrium

	L	R
U	5, 1	0, 0
D	4, 4	1, 5


- If $x = 1$, player 1 plays U, player 2 plays L.
- If $x = 2$, player 1 plays D, player 2 plays L.
- If $x = 3$, player 1 plays D, player 2 plays R.

- We show that no player has an incentive to deviate from the “recommendation” of the mediator:
 - If player 1 gets the recommendation U, he believes player 2 will play L, so his best response is to play U.
 - If player 1 gets the recommendation D, he believes player 2 will play L, R with equal probability, so playing D is a best response.
 - If player 2 gets the recommendation L, he believes player 1 will play U, D with equal probability, so playing L is a best response.
 - If player 2 gets the recommendation R, he believes player 1 will play D, so his best response is to play R.
- Thus the players will follow the mediator’s recommendations.
- With the mediator, the expected payoffs are $(10/3, 10/3)$, strictly higher than what the players could get by randomizing between Nash equilibria.

Correlated equilibrium

- The preceding examples lead us to the notions of correlated strategies and “correlated equilibrium”.
- Let $\Delta(S)$ denote the set of probability measures over the set S . Let R be a random variable taking values in $S = \prod_{i=1}^n S_i$ distributed according to π .
 - An instantiation of R is a pure strategy profile and the i th component of the instantiation will be called the recommendation to player i .
 - Given such a recommendation, player i can use conditional probability to form a posteriori beliefs about the recommendations given to the other players.

- $S_1 = \{A, B\}, S_2 = \{C, D\}, S_3 = \{E, F\}$
- $S = \{(A, C, E), (A, C, F), (A, D, E), (A, D, F), (B, C, E), (B, C, F), (B, D, E), (B, D, F)\}$
- $\Delta(S)$ is a set of probability mass function (PMF) over S
- $\pi \in \Delta(S)$ is a PMF over S
- $R \sim \pi(S)$ is a random variable distributed according to π



Correlated equilibrium

Definition (Correlated equilibrium)

A correlated equilibrium of finite game is a joint probability distribution $\pi \in \Delta(S)$ such that if R is random variable distributed according to π then

$$\sum_{s_{-i} \in S_{-i}} \text{Prob}(R = s | R_i = s_i) u_i(\textcolor{red}{s}_i, s_{-i}) \geq \sum_{s_{-i} \in S_{-i}} \text{Prob}(R = s | R_i = s_i) u_i(\textcolor{blue}{t}_i, s_{-i})$$

For all players i , all $s_i \in S_i$ such that $\text{Prob}(R_i = s_i) > 0$, and all $t_i \in S_i$

- A distribution π is defined to be a correlated equilibrium if no player can ever expect to unilaterally gain by deviating from **his recommendation**, assuming the other players play according to their recommendations.
 - $\textcolor{red}{s}_i$ is a recommendation by R drawn from $\pi \in \Delta(S)$
 - $\textcolor{blue}{t}_i$ is a deviation from this recommendation

Correlated equilibrium

Proposition

A joint probability distribution $\pi \in \Delta(S)$ is a correlated equilibrium of a finite game if and only if

$$\sum_{s_{-i} \in S_{-i}} \pi(s) u_i(s_i, s_{-i}) \geq \sum_{s_{-i} \in S_{-i}} \pi(s) u_i(t_i, s_{-i})$$

For all players i , all $s_i \in S_i$, $t_i \in S_i$ such that $s_i \neq t_i$

Proof:

$$\text{Prob}(R = s | R_i = s_i) = \frac{\pi(s_i, s_{-i})}{\pi(s_i)} = \frac{\pi(s)}{\sum_{t_{-i} \in S_{-i}} \pi(s_i, t_{-i})}$$

$$\sum_{s_{-i} \in S_{-i}} \text{Prob}(R = s | R_i = s_i) u_i(s_i, s_{-i}) \geq \sum_{s_{-i} \in S_{-i}} \text{Prob}(R = s | R_i = s_i) u_i(t_i, s_{-i})$$

$$\sum_{s_{-i} \in S_{-i}} \frac{\pi(s)}{\sum_{t_{-i} \in S_{-i}} \pi(s_i, t_{-i})} u_i(s_i, s_{-i}) \geq \sum_{s_{-i} \in S_{-i}} \frac{\pi(s)}{\sum_{t_{-i} \in S_{-i}} \pi(s_i, t_{-i})} u_i(t_i, s_{-i})$$


- The denominator does not depend on the variable of summation so it can be factored out of the sum and cancelled

Correlated equilibrium

Theorem (Correlated equilibrium)

For every Nash equilibrium σ^* there exists a corresponding correlated equilibrium σ

- Not every correlated equilibrium is equivalent to a Nash equilibrium, (e.g., Battle of Sex game)
 - Correlated equilibrium is **a strictly weaker notion** than Nash



- Any convex combination of the payoffs achievable under correlated equilibria is itself realizable under a correlated equilibrium

Trembling-hand perfect equilibrium

Definition (Trembling-hand perfect equilibrium)

A mixed strategy profile is a (trembling-hand) perfect equilibrium of a normal-form game G if there exists a sequence s^0, s^1, \dots of fully mixed strategy profiles such that $\lim_{n \rightarrow \infty} s^n = s$, and such that for each s^k in the sequence and each player i , the strategy s_i is a best response to the strategies s_{-i}^k .

- Rationalizability is a weaker concept than Nash equilibrium, but perfection is a stronger one
- Perfect equilibria are relevant to one aspect of multiagent learning
- It requires that the equilibrium be robust against slight errors
- One's action out to be the best response not only against the opponent's equilibrium strategies, but also against small perturbation of those.

ϵ – Nash equilibrium

- Players might not care about changing their strategies to a best response when the amount of utility that they could gain by doing so is very small.

Definition (ϵ – Nash equilibrium)

Fix $\epsilon > 0$. A strategy profile $s^* = (s_1^*, \dots, s_n^*)$ is an ϵ -Nash equilibrium if, for all agents i and for all strategies $s_i \neq s_i^*$, $u_i(s_i^*, s_{-i}^*) \geq u_i(s_i, s_{-i}^*) - \epsilon$

	L	R
U	1, 1	0, 0
D	$1 + \frac{\epsilon}{2}, 1$	500, 500

A game with interesting ϵ – Nash equilibrium