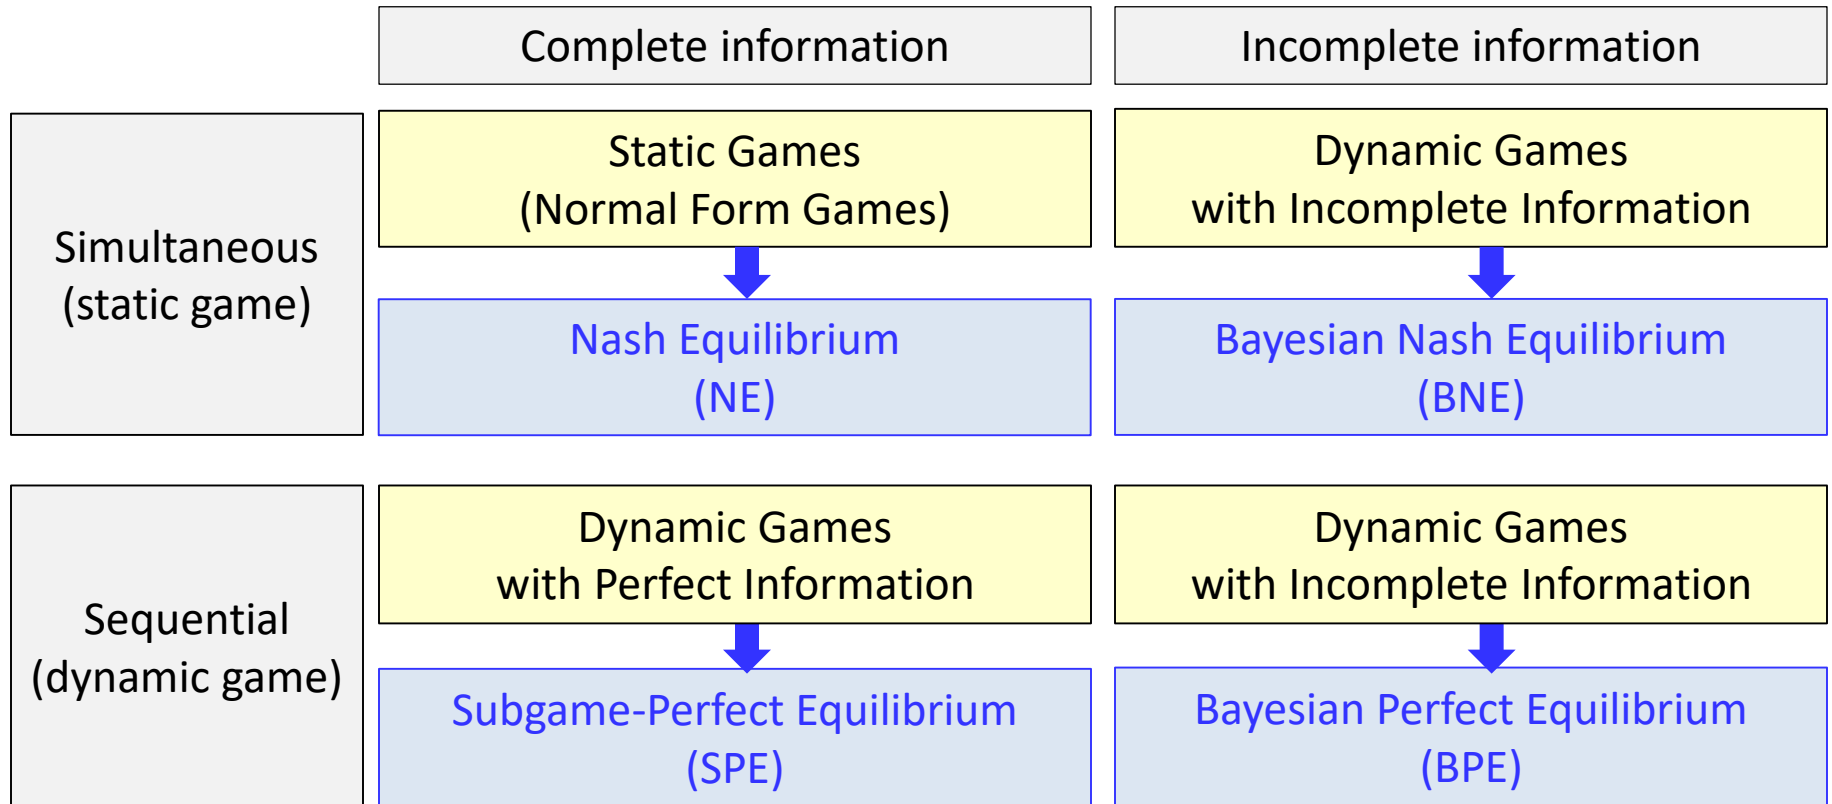


Lecture 10: Bayesian games

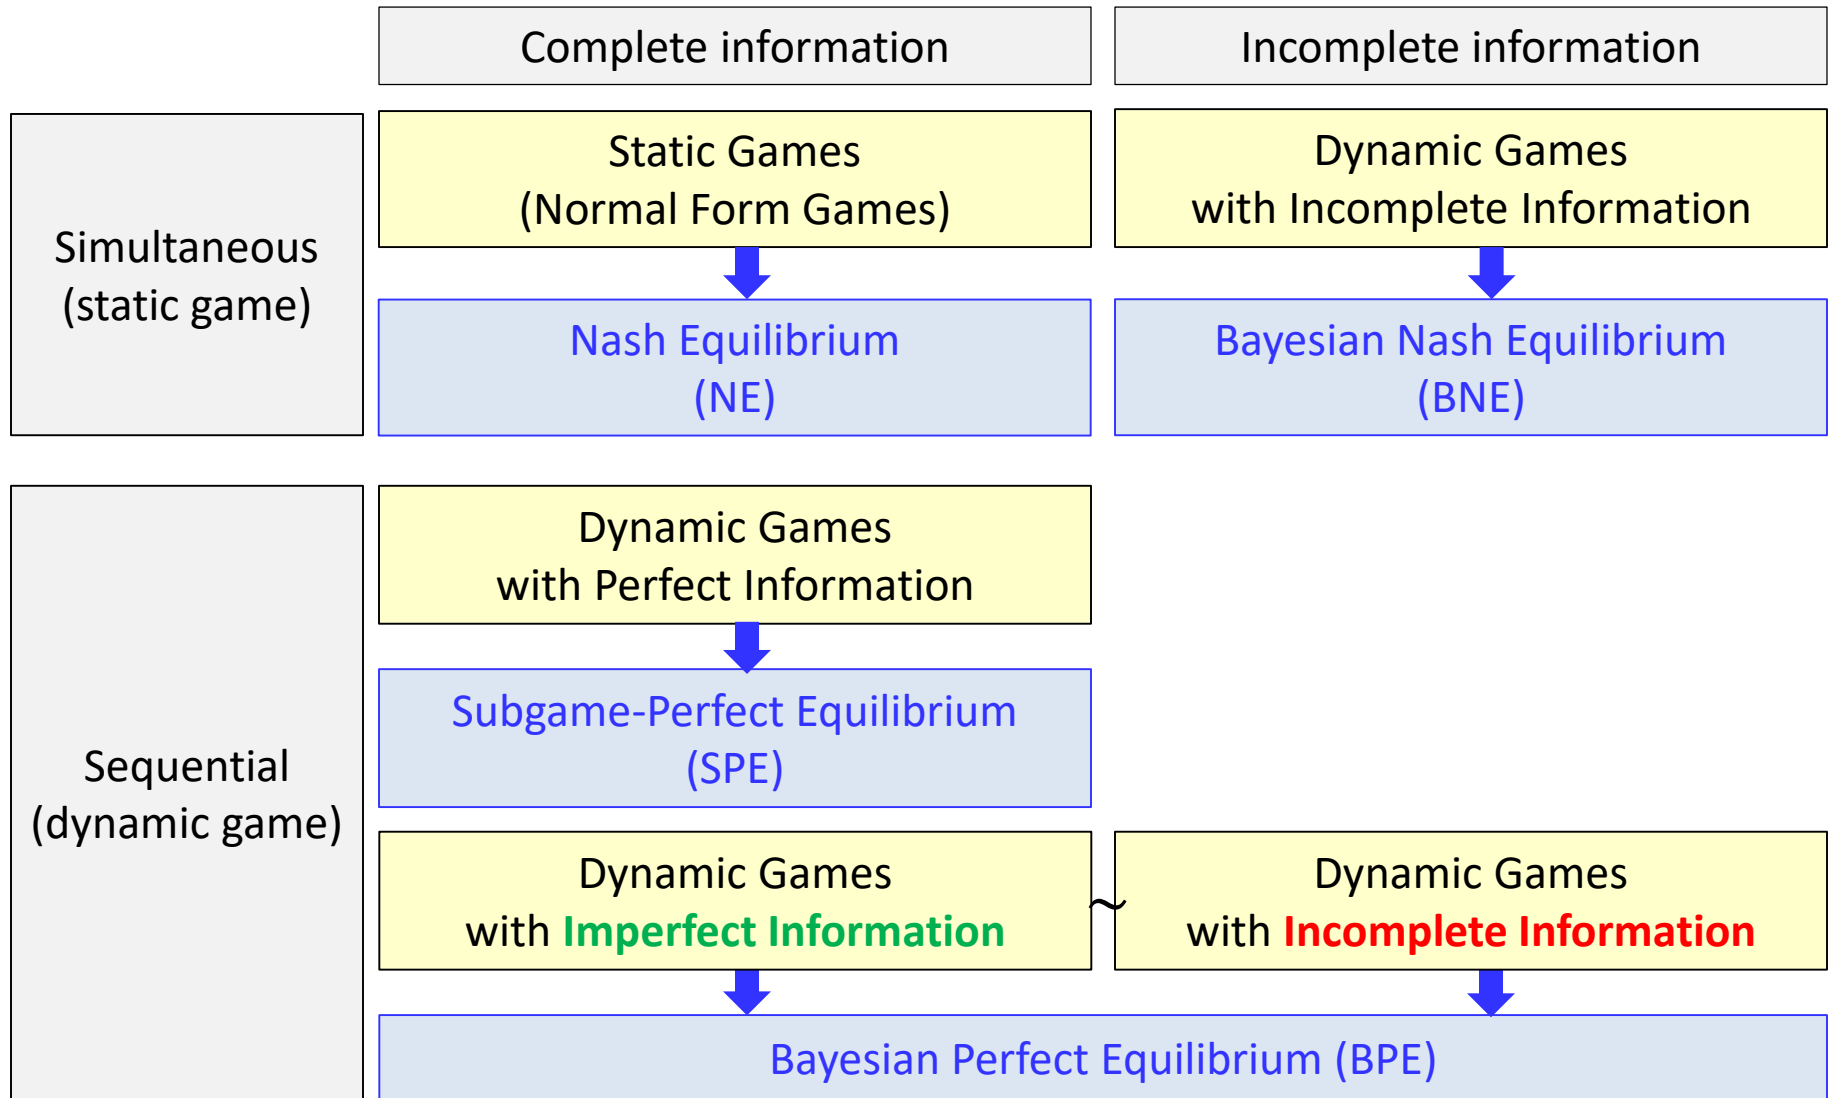
Introduction

- Static games with **complete information**
 - **Normal form games**
 - Dynamic games (sequential games) with **complete information**
 - **Perfect information games**
 - **Imperfect information games**
 - Static games with **incomplete information**
 - **Bayesian games**
 - Dynamic games with **incomplete information**
- **Imperfect Information**: Players do not perfectly observe the actions of other players or forget their own actions.
 - **Incomplete Information**: Players have private information about something relevant to his decision making.
 - Incomplete information introduces **uncertainty** about the game being played
 - Various way to represent the uncertainty about the game (will be discussed)

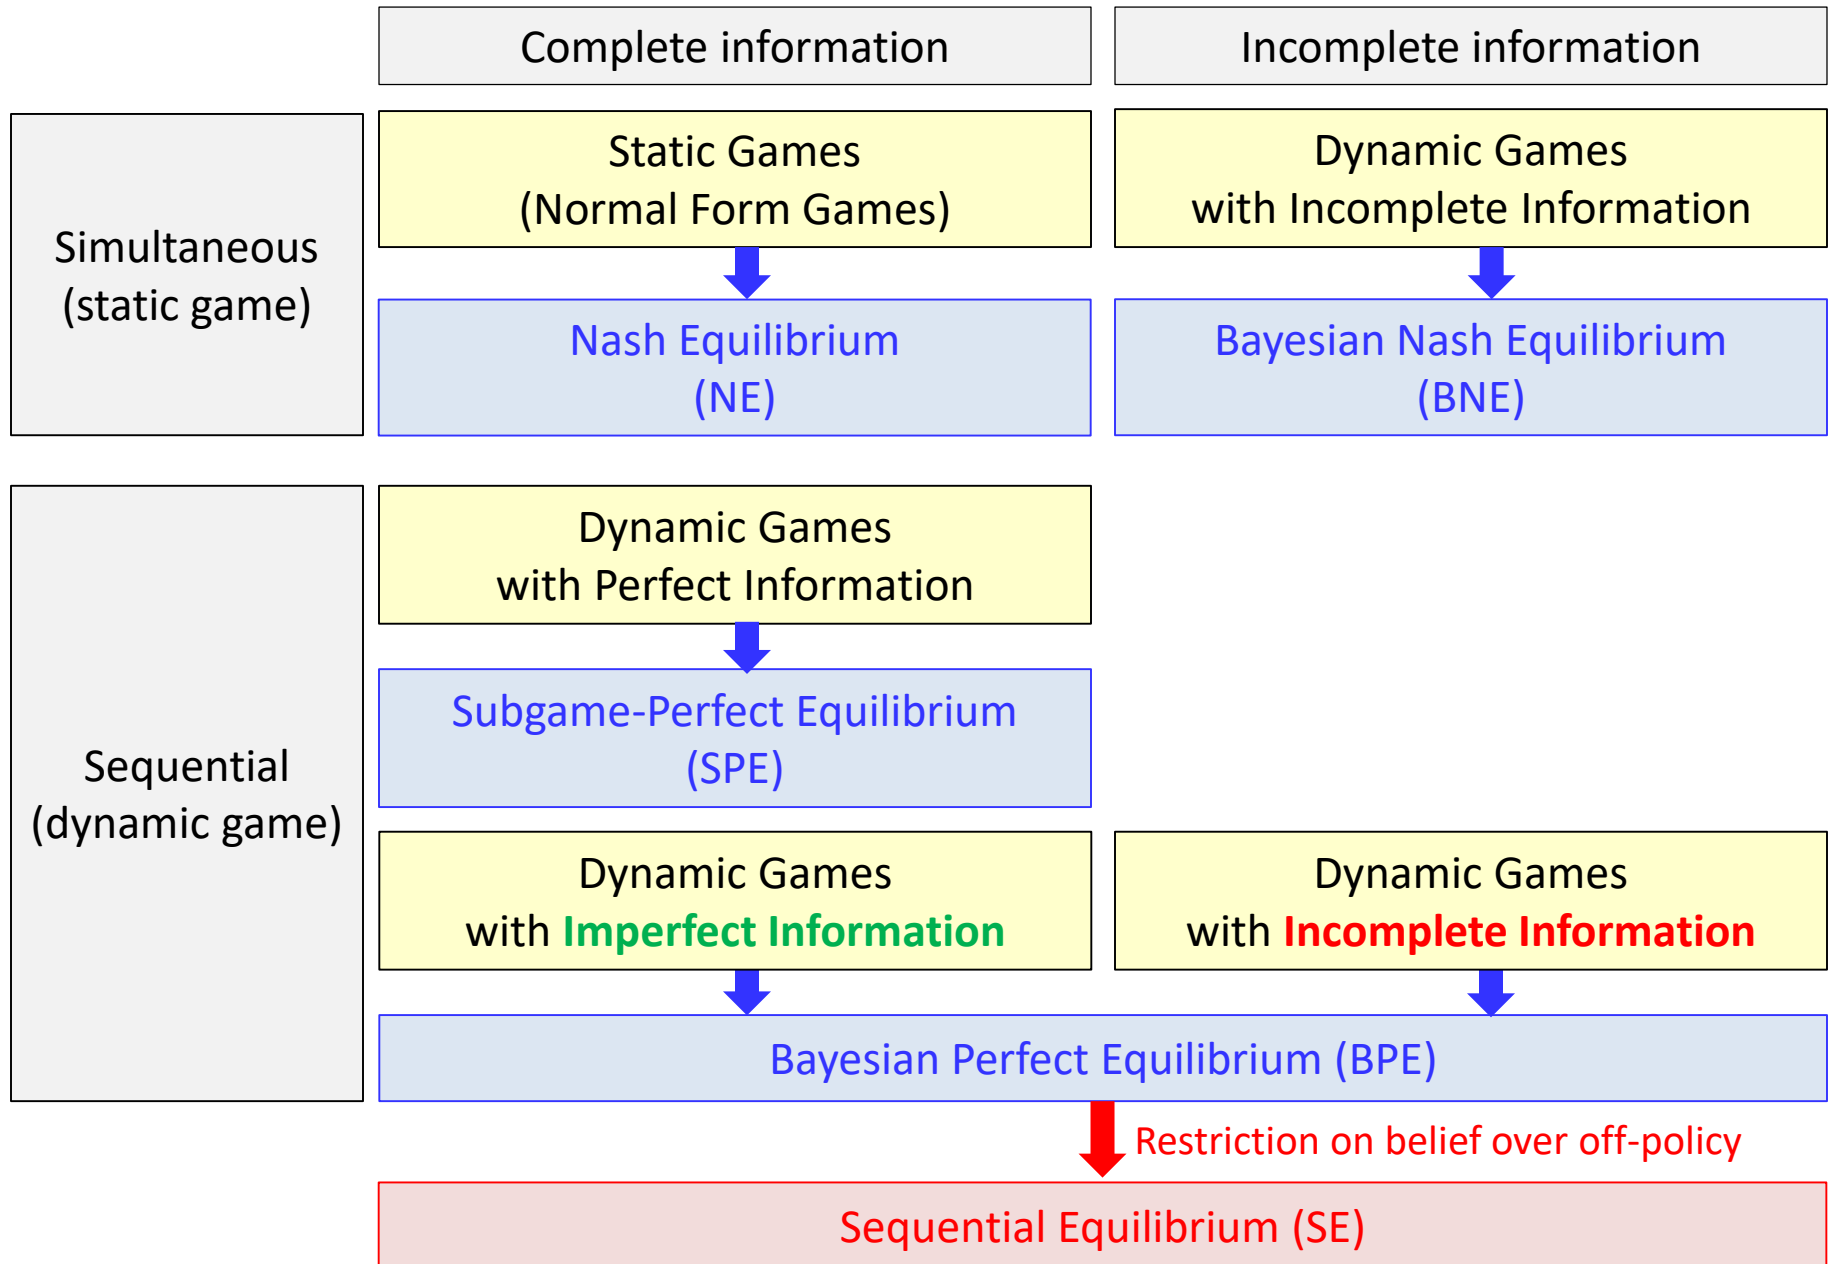
Introduction



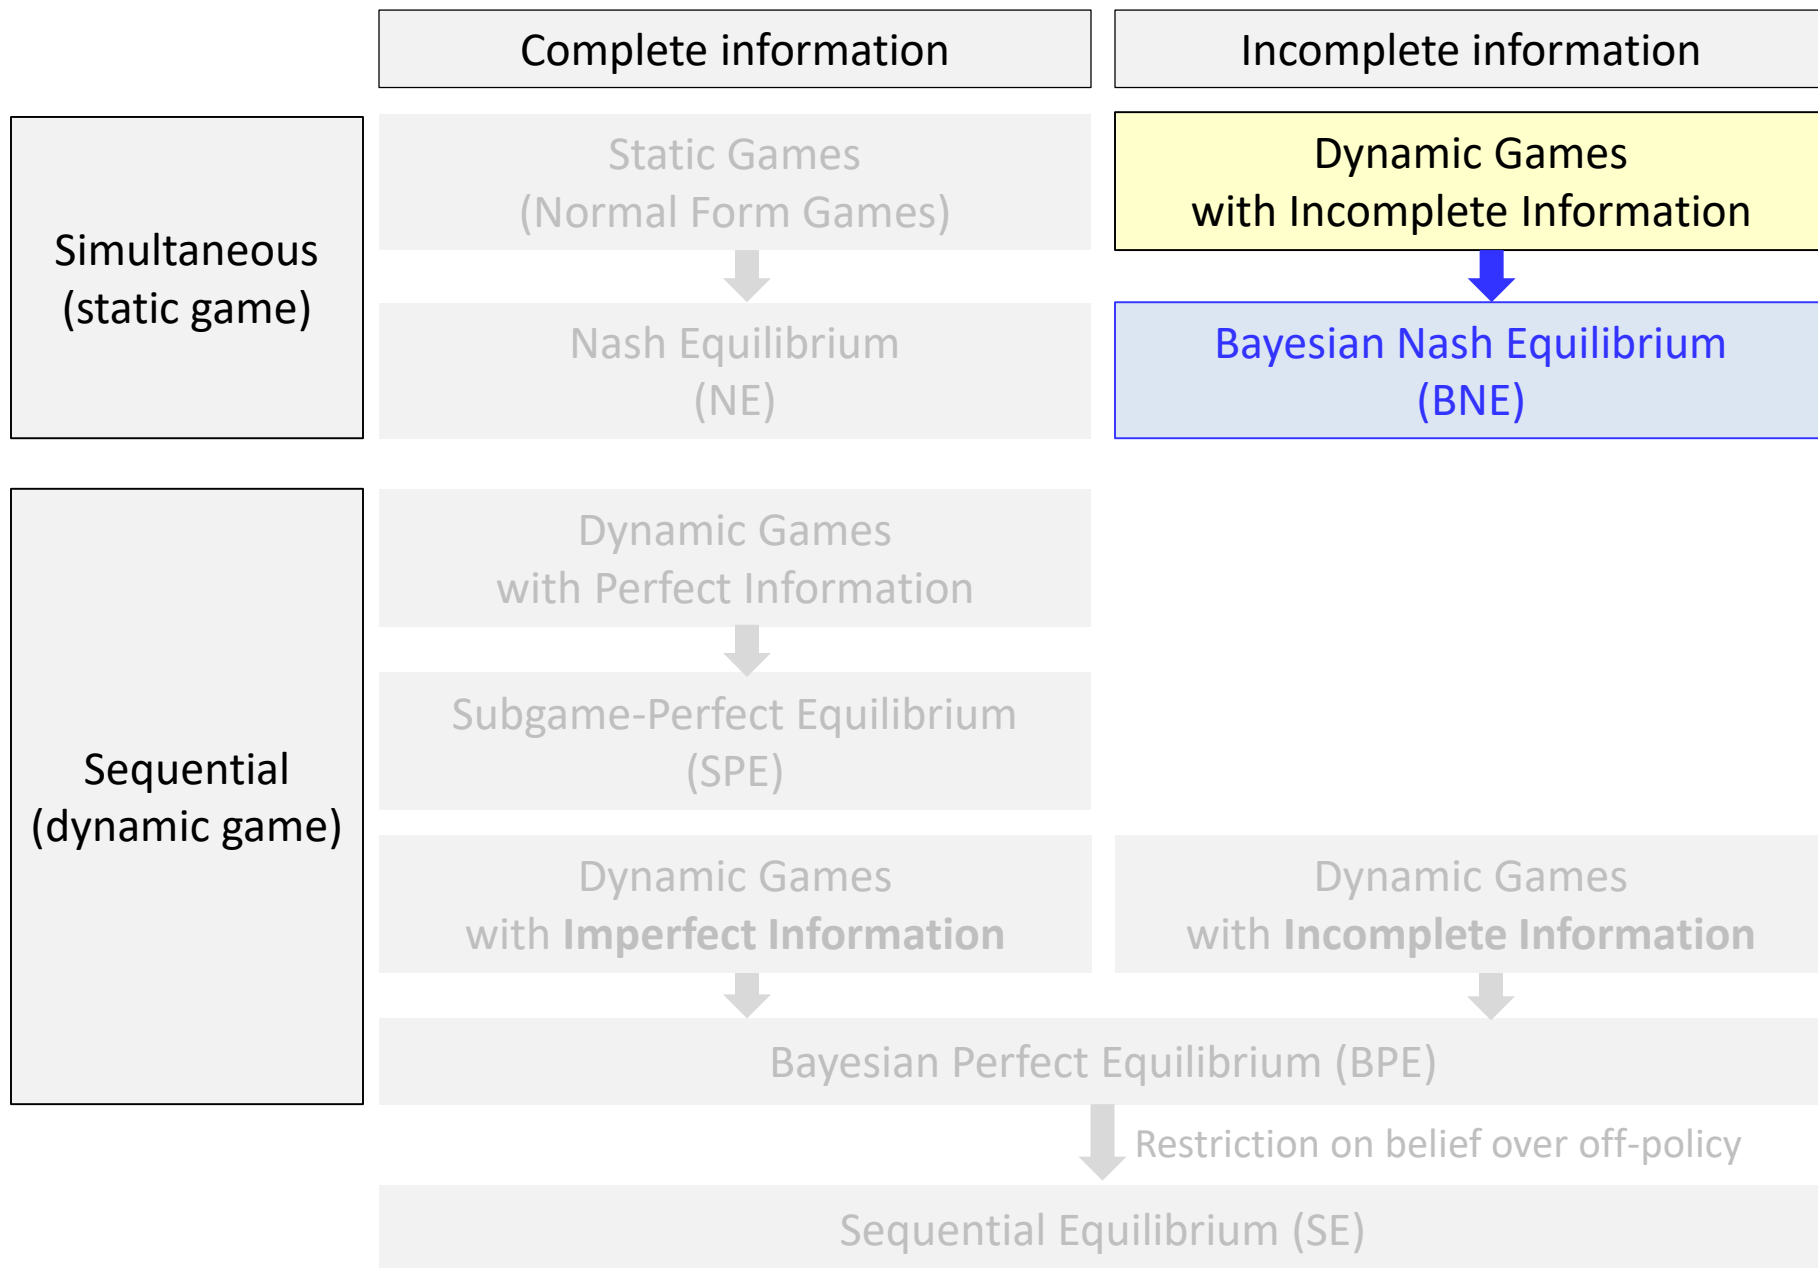
Introduction



Introduction



Introduction



Summary so far....

- So far, we have made an important assumption: The game played is **common knowledge**. That is, the players in the games aware of
 - ✓ Who is playing
 - ✓ What the possible actions of each player are
 - ✓ How outcomes translate into payoffs
- This is true even for imperfect-information games
 - the actual moves of agents are not common knowledge, but the game itself is
- Furthermore, we have assumed that **this knowledge of the game** is itself **common knowledge**
- These assumptions enabled us to lay the methodological foundation for such solution concepts as
 - ✓ Iterated elimination of dominated strategies,
 - ✓ Rationalizability
 - ✓ Nash equilibrium
 - ✓ Subgame-perfect equilibrium
 - ✓ ...

Motivations

- These idealized situations are rarely encountered in reality
- For example, Cournot models of duopolistic competition, each firm may not know
 - ✓ Payoffs of the firms
 - ✓ Actions spaces
 - ✓ Production technologies
 - ✓ Productivity of workers
 - ✓ Cost function of each firm
 - ✓ ...
- Firms have a **reasonably good idea about** their opponents' cost but do not know exactly what they are
- In general, players have some idea about their opponents' characteristics but don't know for sure
- The situation is somewhat similar to **imperfect information game**, in which we analyzed game as follows:
 - Conjecture about the behavior (strategy) of opponents, captured by belief
 - Beliefs and appropriate best Responses should be mutually consistent and correct

Bayesian games

- **Bayesian games** (games of incomplete information) allow us to represent players' **uncertainties about the game**. **We assume**:
 1. All possible games have the same number of agents and the same strategy space for each agent; they differ only in their payoffs.
 2. The beliefs of the different agents are posteriors, obtained by conditioning **a common prior** on individual private signals
- One can imagine many other potential types of uncertainty that players might have about the game
 - how many players are involved
 - what actions are available to each player
 - other aspects of the situation
 - ...
- It turns out that these other types of uncertainty can be reduced to **uncertainty only about payoffs** via problem reformulation

Uncertainty regarding a game

- Imagine that we want to model a situation in which one player is **uncertain about the number of actions available** to the other players
- We can reduce this uncertainty to **uncertainty about payoffs** by padding the game with irrelevant actions

	<i>L</i>	<i>R</i>
<i>U</i>	1, 1	1, 3
<i>D</i>	0, 5	1, 13

	<i>L</i>	<i>C</i>	<i>R</i>
<i>U</i>	1, 1	0, 2	1, 3
<i>D</i>	0, 5	2, 8	1, 13

- Consider the following two-player game, in which the row player does not know whether his opponent has only the two strategies *L* and *R* or also the third one *C*:
 - the newly added column is dominated by the others and will not participate in any Nash equilibrium
 - Indeed, Nash equilibria of the original game are equal to that of the modified game (padded one)

- There are several ways of presenting **Bayesian games**:
 - **Information sets**
 - **Extensive form with chance moves (Nature)**
 - **Epistemic** (i.e., relating to knowledge or to the degree of its validation)

Definition (Bayesian game: information sets)

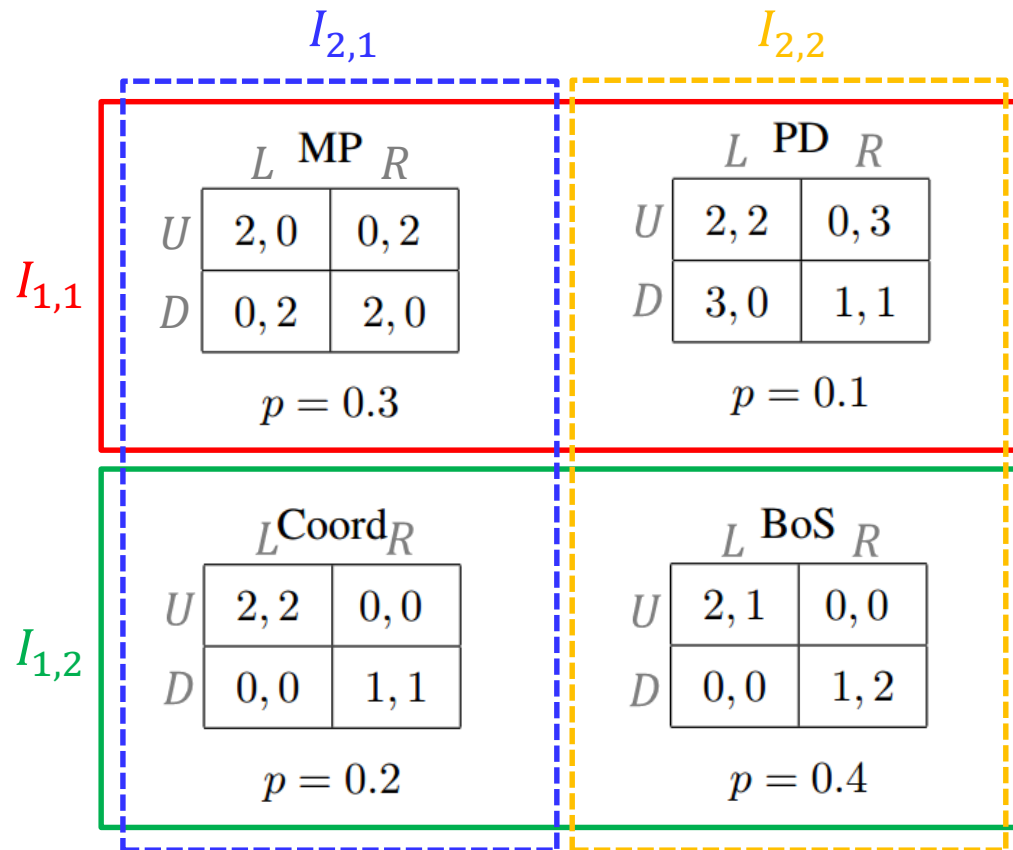
A Bayesian game is a tuple (N, G, P, I) where:

- N is a set of agents
- G is a set of games with n agents each such that if $g, g' \in G$ then for each agent $i \in N$ the strategy space in g is identical to the strategy space in g'
- $P \in \Pi(G)$ is **a common prior over games**, where $\Pi(G)$ is the set of all probability distributions over G ; and
- $I = (I_1, \dots, I_N)$ is a tuple of **partitions of G** , one for each agent

- Equilibrium analysis is possible because,
 - Players have **a common prior over games**
 - Everything about a game must be common knowledge
- Players know their own preferences (information set, types, ...), which in turn will allow us to analyze a players' best response given his assumptions about the behavior of his opponents

Bayesian games represented by information set

Example



- It consists of four 2× 2 games (Matching Pennies, Prisoner's Dilemma, Coordination and Battle of the Sexes), and each agent's partition consists of two equivalence classes

Bayesian games represented in an extensive form game with imperfect information

- A second way of capturing the common prior is to hypothesize a special agent called Nature who makes probabilistic choices

Following **Harsanyi (1967)**, we can change a game of incomplete information into a dynamic game of imperfect information, by making nature as a mover in the game. In such a game, nature chooses player i 's type, but another player j is not perfectly informed about this choice.

- Nature does not have a utility function
- Nature has the unique strategy of randomizing in a commonly known way
- The agents receive individual signals about Nature's choice, and these are captured by their information sets in a standard way
- The information sets capture the fact that agents make their choices without knowing the choices of others

- Thus, we have reduced **games of incomplete information** to **games of imperfect information**, albeit ones with chance moves.

Bayesian games represented in an extensive form game with imperfect information

$I_{2,1}$

$I_{2,2}$

$I_{1,1}$

	L	MP	R
U	2, 0	0, 2	
D	0, 2	2, 0	

$p = 0.3$

	L	PD	R
U	2, 2	0, 3	
D	3, 0	1, 1	

$p = 0.1$

$I_{1,2}$


	L	Coord	R
U	2, 2	0, 0	
D	0, 0	1, 1	

$p = 0.2$

	L	BoS	R
U	2, 1	0, 0	
D	0, 0	1, 2	


$p = 0.4$

- The Bayesian game defined using information set can be represented in extensive form as follows (Harsanyi, Nobel prize 1994)
- Each player knows the distribution of his opponents' type (the common prior assumptions), thus, we can perform analysis



Bayesian games represented in a extensive form game with imperfect information

Entry game example




- **Player 1:** A potential entrant to an industry that has a monopolistic incumbent, player 2
 - Can decide to enter the market (Stay out)
 - Can decide not to enter (Enter)
- **Player 2:** If player 1 enters the market, player 2
 - Can Fight with player 1
 - Can Accommodate with player 1

Extensive form game
Subgame-Perfect equilibrium

Bayesian games represented in an extensive form game with imperfect information

Entry game example



Extensive form game
Subgame-Perfect equilibrium

- **Player 1:** A potential entrant to an industry that has a monopolistic incumbent, player 2
 - Can decide to enter the market (Stay out)
 - Can decide not to enter (Enter)
- **Player 2:** If player 1 enters the market, player 2
 - Can Fight with player 1
 - Can Accommodate with player 1


	F	A
O	<div>$0, 2$</div>	$0, 2$
E	$-1, -1$	<div>$1, 1$</div>

Nash equilibria = $\{(O, F), (E, A)\}$

Subgame perfect equilibrium = (E, A)

Bayesian games represented in an extensive form game with imperfect information


Entry game example



- There are two types of player 2 : $\theta_2 \in \{R, C\}$, and let $P\{\theta_2 = R\} = p$
 - Rational player: Hate fighting
 - Crazy player: Enjoy fighting
- Nature chooses which type of player 2 is playing the game with player 1

Bayesian games represented in an extensive form game with imperfect information


Entry game example



- We assume that players know their own preferences
 - Player 1 is uncertain about player 2's preference (he does not know the type for player 2)
 - But player 2 knows what his preferences are when he needs to make a decision
- If players know their own preferences, but they do not know the preferences of others, What they have to do?
 - Form correct beliefs about the preferences and types of their opponents

Bayesian games represented in an extensive form game with imperfect information

Entry game example



- We need to compute the expected payoff

$$u_1(E, AF) = p \times 1 + (1 - p) \times (-1) = 2p - 1$$


$$u_2(E, AF) = p \times 1 + (1 - p) \times 2 = 2 - p$$

	AA	AF	FA	FF
O	0, 2	0, 2	0, 2	0, 2
E	1, 1	$2p - 1, 2 - p$	$1 - 2p, 1 - 2p$	$-1, 2 - 3p$

- Each player knows the probability distribution over types, and this itself is common knowledge among the players of the game (the common prior assumption)
 - All players know Nature will choose type 1 (rational) with a probability p
- Possible strategies for each player
 - ✓ $s_1 \in \{O, E\}$
 - ✓ $s_2 \in \{AA, AF, FA, FF\}$

Bayesian games represented in a extensive form game with imperfect information

Entry game example



	AA	AF	FA	FF
O	0, 2	0, 2	0, 2	0, 2
E	1, 1	$\frac{1}{3}, \frac{4}{3}$	$-\frac{1}{3}, -\frac{1}{3}$	-1, 0

- There are three pure strategy Nash equilibria
 - $\{(O, FA), (O, FF), (E, AF)\}$
- Only (E, AF) is a subgame-perfect equilibria

Bayesian games represented in a extensive form game with imperfect information

Entry game example

- We modeled this situation as one in which players have uncertainty about the preferences of other players (i.e., uncertainty about the payoffs in the game)
- We assumed that players share the same beliefs about the uncertainty, which allowed us to create a new game for two players with expected payoffs
 - We cannot perform equilibrium analysis without having the common prior
- We have changed the complex and challenging concept of **incomplete information** into a well-known of **imperfect information**, in which
 - Nature choose the player's types
 - We can then use our standard tools of analysis
- Nash equilibrium concept required players to form conjectures, or beliefs, that in equilibrium have to match the choices of their opponents.
- Here we must ask for more
 - **All players agree on the way in which players' types different from each other**
 - **On the way in which Nature chooses among these profiles of types**

Bayesian games represented by epistemic types

- Our third definition uses the notion of an **epistemic type**, or simply a type, as a way of defining uncertainty directly over a **game's utility function**

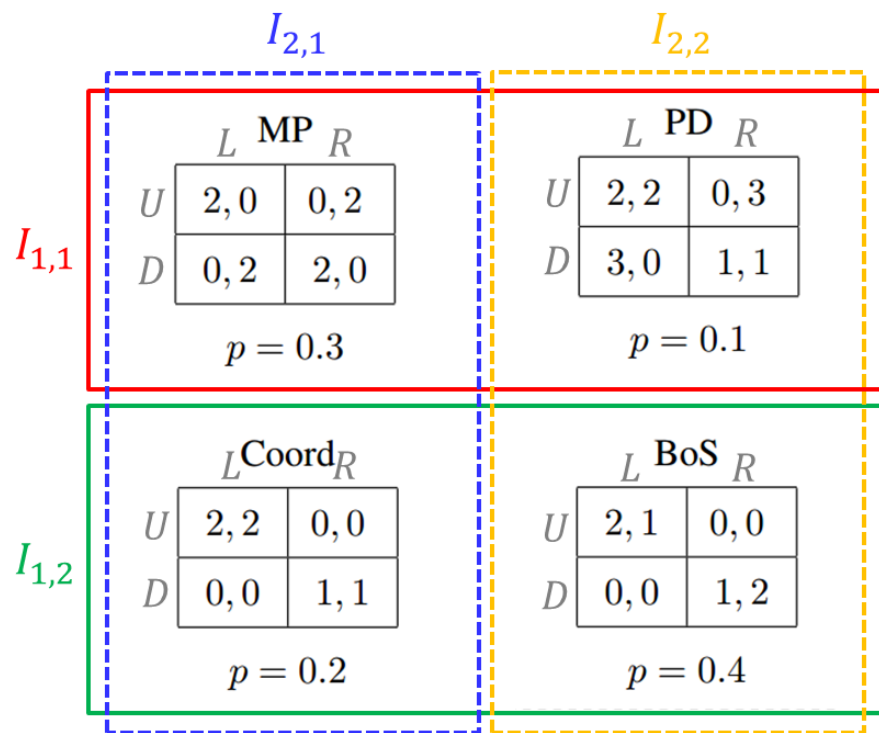
Definition (Bayesian game: epistemic types)

A Bayesian game is a tuple (N, A, Θ, p, u) where:

- N is a set of agents
 - $A = A_1 \times \dots \times A_n$, where A_i is the set of actions available to player i ;
 - $\Theta = \Theta_1 \times \dots \times \Theta_n$, where Θ_i is the **type space** of player i ;
 - $p : \Theta \mapsto [0,1]$ is a **common prior over types**, and
 - $u = (u_1, \dots, u_n)$, where $u_i : A \times \Theta_i \mapsto \mathbb{R}$ is the utility function of player i , which is type dependent,
 - ✓ $u_i(\cdot, \theta)$ is the utility function with a type $\theta \in \Theta$ (**common value setting**)
 - ✓ $u_i(\cdot, \theta_i)$ is the utility function with a type $\theta_i \in \Theta_i$ (**individual value setting**)
- The assumption is that all of the above is common knowledge among the players, and that each agent knows his own type
 - the type of agent **encapsulates** all the information possessed by the agent that is not common knowledge
 - $p(\theta_{-i}|\theta_i)$ is (posterior) conditional distribution on all other types but i given θ_i

Bayesian games represented by epistemic types

Example



a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2



- The joint distribution on these types is as follows:
 - $p(\theta_{1,1}, \theta_{2,1}) = 0.3$; $p(\theta_{1,1}, \theta_{2,2}) = 0.1$; $p(\theta_{1,2}, \theta_{2,1}) = 0.2$; $p(\theta_{1,2}, \theta_{2,2}) = 0.4$

- The **conditional distribution** on player 2 given $\theta_{1,1}$:

- $p(\theta_{2,1} | \theta_{1,1}) = \frac{3}{3+1}$; $p(\theta_{2,2} | \theta_{1,1}) = \frac{1}{3+1}$;

Conjecture about other player's game type given his own (private) type

Bayesian games represented by epistemic types

(static) Bayesian game as a procedure

1. Nature choose a profile of types $\theta = (\theta_1, \theta_2, \dots, \theta_n)$
2. Each player i learns his own type θ_i , which is his private information, and then form posterior beliefs over the other types of players
3. Players simultaneously (hence this is a static game) choose actions $a_i \in A_i$, $i \in N$.
4. Given the players' choices $a = (a_1, a_2, \dots, a_n)$, the payoffs $u_i(a; \theta_i)$ are realized for each player $i \in N$
 - **private values case** : $u_i(a_1, a_2, \dots, a_n; \theta_i)$: private values case
 - ✓ Player's payoff depends only on his type
 - **common values case** : $u_i(a_1, a_2, \dots, a_n; \theta_1, \theta_2, \dots, \theta_n)$
 - ✓ Player's payoff depends on the types of all players

Strategies

- In **imperfect-information extensive-form game**:
 - a pure strategy is a mapping from **every information set** to actions.
- In **a Bayesian game**:
 - A pure strategy $s_i: \Theta_i \mapsto A_i$ is a mapping from **every type** agent i could have to the action he would play if he had that type, i.e., $a_i = s_i(\theta_i)$
 - A mixed strategies in the natural way as probability distributions over pure strategies
 - $s_i(a_i | \theta_i)$ denotes the probability under mixed strategy s_i that agent i plays action a_i , given that i 's type is θ_i .

Expected utility

- Due to uncertainties in Bayesian game, **an expected utility** is used to find an equilibrium
- In a Bayesian game setting, there are three meaningful notions of **expected utility**:
 1. ex post,
 2. ex interim
 3. ex ante.

	his own type	
ex post	Know	Know
ex interim	Know	Don't know
ex ante	Don't know	Don't know

Ex post expected utility

Definition (ex post expected utility)

Agent i 's ex post expected utility in a Bayesian game (N, A, Θ, p, u) , where the agents' strategies are given by s and the agent's types are given by θ , is defined as

$$EU_i(s, \theta) = \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta)$$

Probability of selecting a joint action a given θ

- The types $\theta = (\theta_1, \dots, \theta_n)$ of all agents are given (known)
 - In Bayesian games, no agent will know the others' types
- the only uncertainty concerns the other agents' mixed strategies

$$a \in A = \{(U, L), (U, R), (D, L), (D, R)\}$$

$$\begin{aligned} EU_1(s, (\theta_{1,1}, \theta_{2,1})) &= s_1(U|\theta_{1,1})s_2(L|\theta_{2,1})u_1((U, L), (\theta_{1,1}, \theta_{2,1})) \\ &+ s_1(U|\theta_{1,1})s_2(R|\theta_{2,1})u_1((U, R), (\theta_{1,1}, \theta_{2,1})) \\ &+ s_1(D|\theta_{1,1})s_2(L|\theta_{2,1})u_1((D, L), (\theta_{1,1}, \theta_{2,1})) \\ &+ s_1(D|\theta_{1,1})s_2(R|\theta_{2,1})u_1((D, R), (\theta_{1,1}, \theta_{2,1})) \end{aligned}$$

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Ex interim expected utility

Definition (ex interim expected utility)

Agent i 's ex post expected utility in a Bayesian game (N, A, Θ, p, u) , where i 's type is θ_i and where the agents' strategies are given by the mixed strategy profile s , is defined as

$$EU_i(s, \theta_i) = \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta_{-i}, \theta_i)$$

Or equivalently as

$$EU_i(s, \theta_i) = \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) \underbrace{EU_i(s, (\theta_{-i}, \theta_i))}_{\text{ex post expected utility}}$$

- Agent i must consider **every assignment of types to the other agents θ_{-i}** and **every pure action profile $a \in A$** in order to evaluate his utility function $u_i(a, \theta_{-i}, \theta_i)$
- Because uncertainty over mixed strategies was already handled in the ex post case, we can also write ex interim expected utility as a weighted sum of $EU_i(s, \theta)$ terms.

Ex interim expected utility

$$\begin{aligned}
 EU_i(s, \theta_i) &= \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta_{-i}, \theta_i) \\
 &= \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s, (\theta_{-i}, \theta_i))
 \end{aligned}$$

$$\begin{aligned}
 EU_1(s, \theta_{1,1}) &= p(\theta_{2,1} | \theta_{1,1}) EU_1(s, (\theta_{2,1}, \theta_{1,1})) \\
 &\quad + p(\theta_{2,2} | \theta_{1,1}) EU_1(s, (\theta_{2,2}, \theta_{1,1}))
 \end{aligned}$$

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Ex ante expected utility

Definition (ex ante expected utility)

Agent i 's ex ante expected utility in a Bayesian game (N, A, Θ, p, u) , where the agents' strategies are given by the mixed strategy profile s , is defined as

$$EU_i(s) = \sum_{\theta \in \Theta} p(\theta) \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta)$$

or equivalently as

$$EU_i(s) = \sum_{\theta \in \Theta} p(\theta) \underbrace{EU_i(s, \theta)}_{\text{ex post expected utility}}$$

or again equivalently as

$$EU_i(s) = \sum_{\theta_i \in \Theta_i} p(\theta_i) \underbrace{EU_i(s, \theta_i)}_{\text{ex interim expected utility}}$$

Ex ante expected utility

$$EU_i(s) = \sum_{\theta \in \Theta} p(\theta) \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta)$$

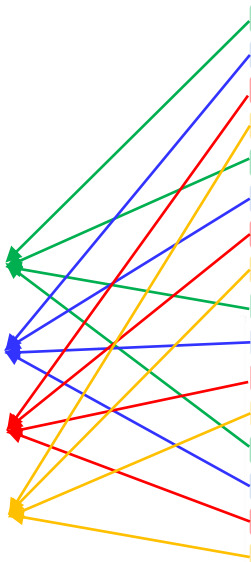
$$= \sum_{\theta \in \Theta} p(\theta) EU_i(s, \theta)$$

$$EU_1(s) = p((\theta_{1,1}, \theta_{2,1})) EU_1(s, (\theta_{2,1}, \theta_{1,1}))$$

$$+ p((\theta_{1,1}, \theta_{2,2})) EU_1(s, (\theta_{1,1}, \theta_{2,2}))$$

$$+ p((\theta_{1,2}, \theta_{2,1})) EU_1(s, (\theta_{1,2}, \theta_{2,1}))$$

$$+ p((\theta_{1,2}, \theta_{2,2})) EU_1(s, (\theta_{1,2}, \theta_{2,2}))$$



a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Bayesian Nash equilibrium

Definition (Best response in a Bayesian game)

The set of agent i 's best responses to mixed-strategy profile s_{-i} are given by

$$BR_i(s_{-i}) = \operatorname{argmax}_{s'_i \in S_i} EU_i(s'_i, s_{-i})$$

- Note that BR_i is a set because there may be many strategies for i that yield the same expected utility

- Because $EU_i(s) = \sum_{\theta_i \in \Theta_i} p(\theta_i) EU_i(s, \theta_i)$

$$EU_i(s, \theta_i) = \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s, (\theta_{-i}, \theta_i))$$

$$EU_i(s, (\theta_{-i}, \theta_i)) = \sum_{a \in A} \left(\prod_{j \in N} s_j(a_j | \theta_j) \right) u_i(a, \theta)$$

Bayesian Nash equilibrium

Definition (Best response in a Bayesian game)

The set of agent i 's best responses to mixed-strategy profile s_{-i} are given by

$$BR_i(s_{-i}) = \operatorname{argmax}_{s'_i \in S_i} EU_i(s'_i, s_{-i})$$

- Note that BR_i is a set because there may be many strategies for i that yield the same expected utility

Why best response is defined in terms of ex-ante expected utility?

- Note that $EU_i(s'_i, s_{-i}) = \sum_{\theta_i \in \Theta_i} p(\theta_i) EU_i(s'_i, s_{-i}, \theta_i)$
- $EU_i(s'_i, s_{-i}, \theta_i)$ does not depend on strategies that i would play if his type were not θ_i
- Thus, we are in fact performing independent maximization of i 's ex interim expected utilities conditioned on each type that he could have
- That is, for each $\theta_i \in \Theta_i$, the following equations should be satisfied

$$EU_i(s_i, s_{-i}, \theta_i) \geq EU_i(s'_i, s_{-i}, \theta_i) \quad \forall s'_i \in S_i$$

Bayesian Nash equilibrium

Definition (Bayes-Nash equilibrium)

A Bayes-Nash equilibrium is a mixed-strategy profile s that satisfies $\forall i \ s_i \in BR_i(s_{-i})$

- This is exactly the same definition we gave for the Nash equilibrium: each agent plays a best response to the strategies of the other players
- Thus, for every player $i \in N$, the following relationship should be satisfied

for each $\theta_i \in \Theta_i$, the following equations should be satisfied

$$EU_i(s_i, s_{-i}, \theta_i) \geq EU_i(s'_i, s_{-i}, \theta_i) \quad \forall s'_i \in S_i$$

for every player $i \in N$, for each $\theta_i \in \Theta_i$

$$\sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s_i, s_{-i}, (\theta_{-i}, \theta_i)) \geq \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s'_i, s_{-i}, (\theta_{-i}, \theta_i)) \quad \forall s'_i \in S_i$$

Bayesian Nash equilibrium

Definition (Bayes-Nash equilibrium)

A Bayes-Nash equilibrium is a mixed-strategy profile s that satisfies $\forall i \ s_i \in BR_i(s_{-i})$

- This is exactly the same definition we gave for the Nash equilibrium: each agent plays a best response to the strategies of the other players
- Thus, for every player $i \in N$, the following relationship should be satisfied

for each $\theta_i \in \Theta_i$, the following equations should be satisfied

$$EU_i(s_i, s_{-i}, \theta_i) \geq EU_i(s'_i, s_{-i}, \theta_i) \quad \forall s'_i \in S_i$$

for every player $i \in N$, for each $\theta_i \in \Theta_i$

$$\sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s_i, s_{-i}, (\theta_{-i}, \theta_i)) \geq \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) EU_i(s'_i, s_{-i}, (\theta_{-i}, \theta_i)) \quad \forall s'_i \in S_i$$

When s_i and s_{-i} are pure strategies



$$\Rightarrow \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) u_i(s_i(\theta_i), s_{-i}(\theta_{-i}); \theta_i) \geq \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) u_i(s'_i(\theta_i), s_{-i}(\theta_{-i}); \theta_i) \quad \forall s'_i \in S_i$$

$$\Rightarrow \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) u_i(s_i(\theta_i), s_{-i}(\theta_{-i}); \theta_i) \geq \sum_{\theta_{-i} \in \Theta_{-i}} p(\theta_{-i} | \theta_i) u_i(a_i, s_{-i}(\theta_{-i}); \theta_i) \quad \forall a_i \in A_i$$

Computing Bayesian Nash equilibrium

- As we did with extensive-form games, we can construct a normal-form representation that corresponds to a given Bayesian game.
- Each agent i 's payoff given a pure-strategy profile s is his **ex ante expected utility** under s
- The Bayes–Nash equilibria of a Bayesian game are precisely the Nash equilibria of its induced normal form
- This fact allows us to note that Nash's theorem applies directly to Bayesian games, and hence that Bayes–Nash equilibria always exist.

Computing Bayesian Nash equilibrium

$I_{2,1}$

$I_{2,2}$

$I_{1,1}$

	L	MP	R
U	2, 0		0, 2
D	0, 2		2, 0

$p = 0.3$

	L	PD	R
U	2, 2		0, 3
D	3, 0		1, 1

$p = 0.1$

$I_{1,2}$

	L	Coord	R
U	2, 2		0, 0
D	0, 0		1, 1

$p = 0.2$

	L	BoS	R
U	2, 1		0, 0
D	0, 0		1, 2

$p = 0.4$

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

- Each agent has four possible pure strategies (two types and two actions).
 - Give player 1's type $\theta_{1,1}$: we can select U or D
 - Give player 1's type $\theta_{1,2}$: we can select U or D $\left. \begin{array}{l} \text{Give player 1's type } \theta_{1,1}: \text{ we can select } U \text{ or } D \\ \text{Give player 1's type } \theta_{1,2}: \text{ we can select } U \text{ or } D \end{array} \right\} s_1 = (s_1(\theta_{1,1}), s_1(\theta_{1,2}))$
- Then player 1's four strategies in the Bayesian game can be labeled UU, UD, DU , and DD
- Similarly, we can denote the strategies of player 2 by RR, RL, LR , and LL .

Computing Bayesian Nash equilibrium

- We can convert the Bayesian game into an induced normal form game, each payoff is ex ante expected value

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Bayesian game



	LL	LR	RL	RR
UU	2, 1	1, 0.7	1, 1.2	0, 0.9
UD	0.8, 0.2	1, 1.1	0.4, 1	0.6, 1.9
DU	1.5, 1.4	0.5, 1.1	1.7, 0.4	0.7, 0.1
DD	0.3, 0.6	0.5, 1.5	1.1, 0.2	1.3, 1.1

Induced normal form game

Computing Bayesian Nash equilibrium

- Given a particular signal, the agent can compute the posterior probabilities and re-compute the expected utility of any given strategy vector
- Once the row agent gets the signal $\theta_{1,1}$ he can update the expected payoffs and compute the new game shown bellow

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Bayesian game



Bayesian
updating

	LL	LR	RL	RR
UU	2, 0.5	1.5, 0.75	0.5, 2	0, 2.25
UD	2, 0.5	1.5, 0.75	0.5, 2	0, 2.25
DU	0.75, 1.5	0.25, 1.75	2.25, 0	1.75, 0.25
DD	0.75, 1.5	0.25, 1.75	2.25, 0	1.75, 0.25

Induced normal form game

Computing Bayesian Nash equilibrium

- Given a particular signal, the agent can compute the posterior probabilities and re-compute the expected utility of any given strategy vector
- Once the row agent gets the signal $\theta_{1,1}$ he can update the expected payoffs and compute the new game shown bellow

a_1	a_2	θ_1	θ_2	u_1	u_2
U	L	$\theta_{1,1}$	$\theta_{2,1}$	2	0
U	L	$\theta_{1,1}$	$\theta_{2,2}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,1}$	2	2
U	L	$\theta_{1,2}$	$\theta_{2,2}$	2	1
U	R	$\theta_{1,1}$	$\theta_{2,1}$	0	2
U	R	$\theta_{1,1}$	$\theta_{2,2}$	0	3
U	R	$\theta_{1,2}$	$\theta_{2,1}$	0	0
U	R	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	L	$\theta_{1,1}$	$\theta_{2,1}$	0	2
D	L	$\theta_{1,1}$	$\theta_{2,2}$	3	0
D	L	$\theta_{1,2}$	$\theta_{2,1}$	0	0
D	L	$\theta_{1,2}$	$\theta_{2,2}$	0	0
D	R	$\theta_{1,1}$	$\theta_{2,1}$	2	0
D	R	$\theta_{1,1}$	$\theta_{2,2}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,1}$	1	1
D	R	$\theta_{1,2}$	$\theta_{2,2}$	1	2

Bayesian game



Bayesian
updating

	LL	LR	RL	RR
UU	2, 0.5	1.5, 0.75	0.5, 2	0, 2.25
UD	2, 0.5	1.5, 0.75	0.5, 2	0, 2.25
DU	0.75, 1.5	0.25, 1.75	2.25, 0	1.75, 0.25
DD	0.75, 1.5	0.25, 1.75	2.25, 0	1.75, 0.25

Induced normal form game

- The induced normal form game can be used to find the best response for player 1 given his signal $\theta_{1,1}$
- However, it cannot be used for Bayesian Nash equilibrium, because the signal is not common knowledge

Example: Inefficient trade and adverse selection

- Market allocate goods to the people who value them the most
 - If a good is misallocated so that some people who have it value it less than people who do not
 - Market pressures will cause the price of that good to increase until the current owners will prefer to sell it rather than hold on to it
 - The people who value it more will be willing to pay that price
- It is important to understand the extent to which these arguments stand or fall in the face of incomplete information, when some people are better informed about the value of goods than others

Example: Inefficient trade and adverse selection

- Imagine a scenario in which player 1 owns an orange grove
 - The yields of fruit depends on the quality of the soil and other local conditions
 - There are three types of land representing soil quality $\theta_1 \in \Theta_1 = \{L, M, H\}$
 - Only player 1 knows the quality of the land
 - The value for player 1 of owning land of type θ_1 is given as


$$v_1(\theta_1) = \begin{cases} 10 & \text{if } \theta_1 = L \\ 20 & \text{if } \theta_1 = M \\ 30 & \text{if } \theta_1 = H \end{cases}$$

- Imagine that player 2 is a soybean grower who is considering the purchase of this land for production
- Player 2's monetary-equivalent values are given by

$$v_2(\theta_1) = \begin{cases} 14 & \text{if } \theta_1 = L \\ 24 & \text{if } \theta_1 = M \\ 34 & \text{if } \theta_1 = H \end{cases}$$


- Player 2 knows only that the quality is distributed equally among the three options
- i.e., $P(\theta_1 = L) = P(\theta_1 = M) = P(\theta_1 = H) = 1/3$

Example: Inefficient trade and adverse selection




- Consider the following game
 - Player 2 makes a take-it-or-leave-it price offer to player 1
 - Player 1 accept (A) or reject (R) the offer
 - The game ends with either a transfer of land for the suggested price or no transfer
- A strategy for player 2 is a single price offer $p \geq 0$
- A Pure strategy for player 1 is a mapping from the offered price and his type space Θ_1 to a response, $s_1: [0, \infty) \times \Theta_1 \rightarrow \{A, R\}$

Example: Inefficient trade and adverse selection



- Trade can occur in a Bayesian Nash equilibrium only if it involves the lowest type of player 1 trading. Furthermore any price $p^* \in [10, 14]$ can be supported as a Bayesian Nash equilibrium
- Why?

Example: Inefficient trade and adverse selection



- Trade can occur in a Bayesian Nash equilibrium only if it involves the lowest type of player 1 trading. Furthermore any price $p^* \in [10, 14]$ can be supported as a Bayesian Nash equilibrium

- Why?

$$s_1(\theta_1) = \begin{cases} A & \text{if and only if } p \geq p^* \text{ and otherwise when } \theta_1 = L \\ A & \text{if and only if } p \geq 20 \text{ and otherwise when } \theta_1 = M \\ A & \text{if and only if } p \geq 30 \text{ and otherwise when } \theta_1 = H \end{cases}$$